

⚙ Execution Log Summary:

```
python MRTotalRevenueByStore.py \  
> -r hadoop \  
> hdfs:///user/hadoop/liquor_sales/Liquor_Sales_Final.csv \  
> --output-dir hdfs:///user/hadoop/output/total_revenue_by_store/  
No configs found; falling back on auto-configuration  
No configs specified for hadoop runner  
Looking for hadoop binary in $PATH...  
Found hadoop binary: /usr/bin/hadoop  
Using Hadoop version 2.8.5  
Looking for Hadoop streaming jar in /home/hadoop/contrib...  
Looking for Hadoop streaming jar in /usr/lib/hadoop-mapreduce...  
Found Hadoop streaming jar: /usr/lib/hadoop-mapreduce/hadoop-streaming.jar  
Creating temp directory  
/tmp/MRTotalRevenueByStore.hadoop.20250428.160953.141935  
uploading working dir files to  
hdfs:///user/hadoop/tmp/mrjob/MRTotalRevenueByStore.hadoop.20250428.160953.14  
1935/files  
/wd...  
Copying other local files to  
hdfs:///user/hadoop/tmp/mrjob/MRTotalRevenueByStore.hadoop.20250428.160953.14  
1935/files  
/  
Running step 1 of 1...  
packageJobJar: [] [/usr/lib/hadoop/hadoop-streaming-2.8.5-amzn-6.jar]  
/tmp/streamjob8813397535975255142.jar tmpDir=null  
Connecting to ResourceManager at ip-172-31-99-99.ec2.internal/172.31.99.99:8032  
Connecting to ResourceManager at ip-172-31-99-99.ec2.internal/172.31.99.99:8032  
Loaded native gpl library  
Successfully loaded & initialized native-lzo library [hadoop-lzo rev  
ff8f5709577defb6b78cdc1f98cfe129c4b6fe46]  
Total input files to process : 1  
number of splits:36  
Submitting tokens for job: job_1745854108616_0001  
Submitted application application_1745854108616_0001  
The url to track the job:  
http://ip-172-31-99-99.ec2.internal:20888/proxy/application_1745854108616_0001/  
Running job: job_1745854108616_0001  
Job job_1745854108616_0001 running in uber mode : false  
map 0% reduce 0%  
map 4% reduce 0%  
map 8% reduce 0%
```

map 10% reduce 0%
map 13% reduce 0%
map 17% reduce 0%
map 20% reduce 0%
map 21% reduce 0%
map 23% reduce 0%
map 25% reduce 0%
map 26% reduce 0%
map 28% reduce 0%
map 30% reduce 0%
map 31% reduce 0%
map 32% reduce 0%
map 33% reduce 0%
map 34% reduce 0%
map 36% reduce 0%
map 38% reduce 0%
map 39% reduce 0%
map 40% reduce 0%
map 42% reduce 0%
map 43% reduce 0%
map 44% reduce 2%
map 45% reduce 2%
map 46% reduce 2%
map 49% reduce 2%
map 51% reduce 2%
map 52% reduce 2%
map 54% reduce 2%
map 55% reduce 2%
map 56% reduce 2%
map 58% reduce 2%
map 59% reduce 2%
map 61% reduce 2%
map 61% reduce 5%
map 67% reduce 6%
map 68% reduce 6%
map 69% reduce 6%
map 73% reduce 6%
map 76% reduce 6%
map 77% reduce 6%
map 77% reduce 7%
map 78% reduce 7%
map 79% reduce 7%
map 80% reduce 7%
map 81% reduce 7%
map 83% reduce 7%

map 83% reduce 8%
map 89% reduce 8%
map 92% reduce 8%
map 92% reduce 9%
map 94% reduce 9%
map 97% reduce 9%
map 100% reduce 9%
map 100% reduce 14%
map 100% reduce 19%
map 100% reduce 30%
map 100% reduce 35%
map 100% reduce 42%
map 100% reduce 44%
map 100% reduce 46%
map 100% reduce 48%
map 100% reduce 50%
map 100% reduce 51%
map 100% reduce 54%
map 100% reduce 66%
map 100% reduce 68%
map 100% reduce 69%
map 100% reduce 71%
map 100% reduce 85%
map 100% reduce 99%
map 100% reduce 100%

Job job_1745854108616_0001 completed successfully

Output directory: hdfs:///user/hadoop/output/total_revenue_by_store/

Counters: 51

File Input Format Counters

Bytes Read=4724934886

File Output Format Counters

Bytes Written=36760

File System Counters

FILE: Number of bytes read=49280181

FILE: Number of bytes written=107288009

FILE: Number of large read operations=0

FILE: Number of read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=4724939998

HDFS: Number of bytes written=36760

HDFS: Number of large read operations=0

HDFS: Number of read operations=129

HDFS: Number of write operations=14

Job Counters

Data-local map tasks=36

Killed map tasks=1
Killed reduce tasks=1
Launched map tasks=36
Launched reduce tasks=8
Total megabyte-milliseconds taken by all map tasks=1448712192
Total megabyte-milliseconds taken by all reduce tasks=991340544
Total time spent by all map tasks (ms)=943172
Total time spent by all maps in occupied slots (ms)=1448712192
Total time spent by all reduce tasks (ms)=322702
Total time spent by all reduces in occupied slots (ms)=991340544
Total vcore-milliseconds taken by all map tasks=943172
Total vcore-milliseconds taken by all reduce tasks=322702
Map-Reduce Framework
CPU time spent (ms)=631410
Combine input records=0
Combine output records=0
Failed Shuffles=0
GC time elapsed (ms)=7104
Input split bytes=5112
Map input records=19485373
Map output bytes=221833789
Map output materialized bytes=50500449
Map output records=19485372
Merged Map outputs=252
Physical memory (bytes) snapshot=18596872192
Reduce input groups=2480
Reduce input records=19485372
Reduce output records=2480
Reduce shuffle bytes=50500449
Shuffled Maps =252
Spilled Records=38970744
Total committed heap usage (bytes)=18113626112
Virtual memory (bytes) snapshot=151015665664
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
job output is in hdfs:///user/hadoop/output/total_revenue_by_store/
Removing HDFS temp directory
hdfs:///user/hadoop/tmp/mrjob/MRTotalRevenueByStore.hadoop.20250428.160953.14
1935...

Removing temp directory

/tmp/MRTotalRevenueByStore.hadoop.20250428.160953.141935...

[hadoop@ip-172-31-99-99 ~]\$